



BATCH PIR AND LABELED PSI WITH OBLIVIOUS CIPHERTEXT COMPRESSION

ALEXANDER BIENSTOCK,¹ SARVAR PATEL,² JOON
YOUNG SEO,² KEVIN YEO³


¹JPMORGAN AI RESEARCH & ALGOCRYPT COE (WORK DONE WHILE AT NYU AND
INTERNING AT GOOGLE)

²GOOGLE

³GOOGLE AND COLUMBIA UNIVERSITY

PRIVATE INFORMATION RETRIEVAL (PIR)

- Client has single item x , server has $Y = \{(y_1, v_1), \dots, (y_n, v_n)\}$



(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

PRIVATE INFORMATION RETRIEVAL (PIR)

- Client has single item x , server has $Y = \{(y_1, v_1), \dots, (y_n, v_n)\}$
 - Client wants to privately check if $x = y_i$ in Y



$Req(x)$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

PRIVATE INFORMATION RETRIEVAL (PIR)

- Client has single item x , server has $Y = \{(y_1, v_1), \dots, (y_n, v_n)\}$
 - Client wants to privately check if $x = y_i$ in Y



$Resp(x)$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

PRIVATE INFORMATION RETRIEVAL (PIR)

- Client has single item x , server has $Y = \{(y_1, v_1), \dots, (y_n, v_n)\}$
 - Client wants to privately check if $x = y_i$ in Y
 - If so retrieve v_i



v_i if $x = y_i$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

PRIVATE INFORMATION RETRIEVAL (PIR)

- Client has single item x , server has $Y = \{(y_1, v_1), \dots, (y_n, v_n)\}$
 - Client wants to privately check if $x = y_i$ in Y
 - If so retrieve v_i
 - “Privately” – the server should not learn what x is



v_i if $x = y_i$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

PRIVATE INFORMATION RETRIEVAL (PIR)

- Client has single item x , server has $Y = \{(y_1, v_1), \dots, (y_n, v_n)\}$
 - Client wants to privately check if $x = y_i$ in Y
 - If so retrieve v_i
 - “Privately” – the server should not learn what x is
- This talk: no preprocessing



v_i if $x = y_i$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

BATCH PRIVATE INFORMATION RETRIEVAL (PIR)

- *Batch* Private Information Retrieval (Batch PIR)
 - Client has $X = \{x_1, \dots, x_\ell\}$, server has $Y = \{(y_1, v_1), \dots, (y_n, v_n)\}$
 - Client wants to privately check if each $x_j = y_i$ in Y ; if so retrieve v_i
 - “Privately” – the server should not learn what X is




$$X = \{x_1, \dots, x_\ell\}$$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

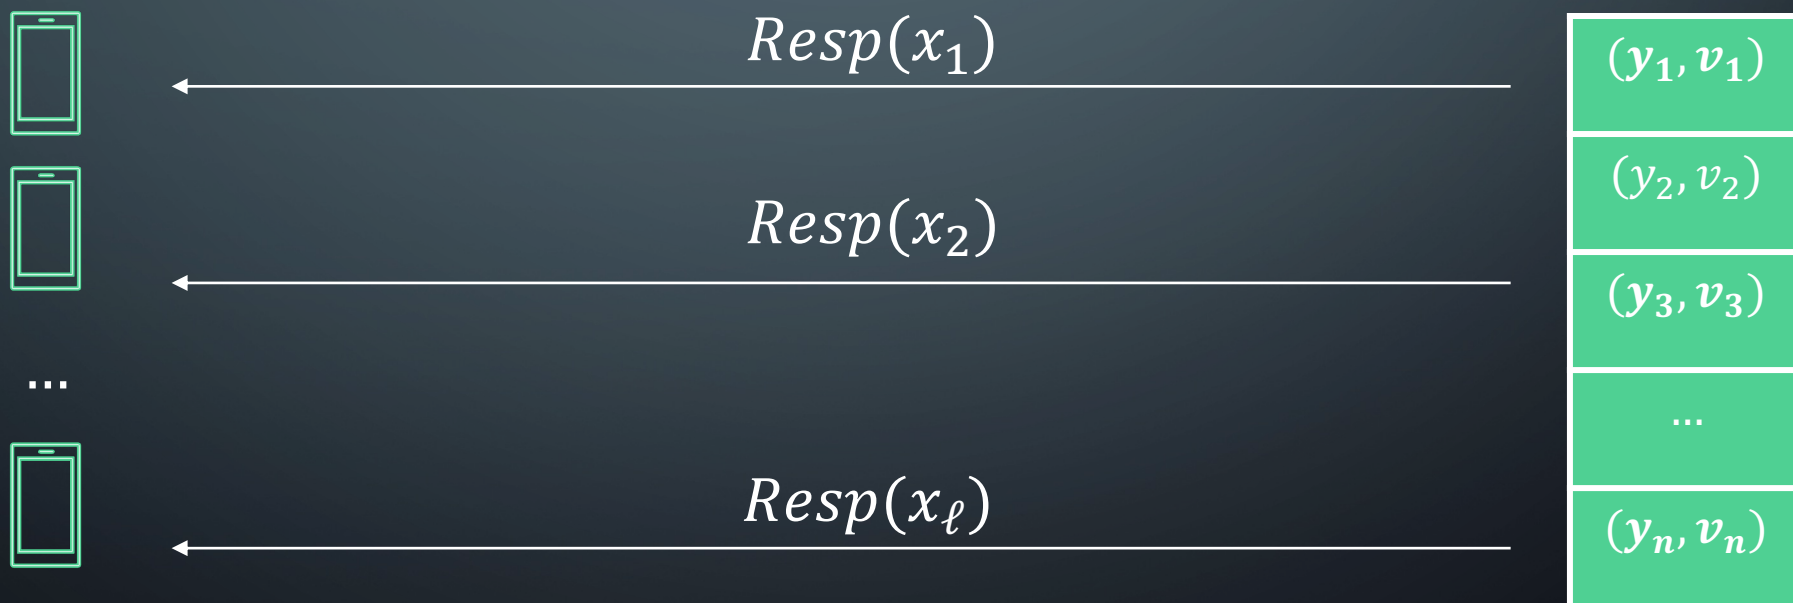
NAÏVE BATCH PIR FROM SINGLE REQUEST PIR

- Run single request PIR ℓ times






NAÏVE BATCH PIR FROM SINGLE REQUEST PIR

- Run single request PIR ℓ times



NAÏVE BATCH PIR FROM SINGLE REQUEST PIR




- Run single request PIR ℓ times

 v_{i_1} if $x_1 = y_{i_1}$
 v_{i_2} if $x_2 = y_{i_2}$
...
 v_{i_ℓ} if $x_\ell = y_{i_\ell}$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

NAÏVE BATCH PIR FROM SINGLE REQUEST PIR

- Run single request PIR ℓ times
 - Bad: factor ℓ blowup in overhead (server comp is now $O(\ell n)$)

 v_{i_1} if $x_1 = y_{i_1}$
 v_{i_2} if $x_2 = y_{i_2}$
...
 v_{i_ℓ} if $x_\ell = y_{i_\ell}$

(y_1, v_1)
(y_2, v_2)
(y_3, v_3)
...
(y_n, v_n)

EFFICIENT BATCH PIR

- [ACLS18] use “Cuckoo Hashing” to remove factor ℓ server blowup
- But communication increases $1.5\times$
 - 1.5ℓ (encrypted) PIR requests/responses; 0.5ℓ of them “useless”
- Can we get rid of communication overhead?
 - This work: Yes – Oblivious Ciphertext (De)compression

OBLIVIOUS CIPHERTEXT COMPRESSION

PIR Responses

$\mathbf{c} =$

$Enc(20)$	$Enc(0)$	$Enc(0)$	$Enc(14)$	$Enc(3)$	$Enc(0)$	$Enc(35)$
-----------	----------	----------	-----------	----------	----------	-----------

ℓ -out-of- n non-zero

- Server compresses \mathbf{c} to \mathbf{c}' of size $\ell + \epsilon$, ϵ small

$\mathbf{c}' =$

--	--	--	--	--

- Client decompresses and decrypts \mathbf{c}' to $\mathbf{m} = (m_1, \dots, m_\ell)$ (nonzeros of \mathbf{c})

$\mathbf{m} =$

20	14	3	35
----	----	---	----

Server must
not learn non-
zero m_i 's or
their locations

OUR MAIN TOOL: RANDOM-BAND MATRICES [DW19]

- Given $w \leq m$, sample $m \times \ell$ matrix M column-by-column as follows:
 - Pick random starting location from $1, \dots, m - w$
 - Pick random w -bit string
 - Embed random w -bit string at starting location; everything else 0



OUR MAIN TOOL: RANDOM-BAND MATRICES [DW19]

- **Theorem** [DW19]: There exists constant $\epsilon > 0$ such that if $m = (1 + \epsilon) \cdot \ell$, $w = O((\log \ell)/\epsilon)$, then randomly sampled $m \times \ell$ random band matrices have full column rank, time $O(\ell w)$ Gaussian elimination, and time $O(\ell w)$ vector multiplication, with **high probability**
- **Our Theorem:** There exists constant $\epsilon > 0$ such that if $m = (1 + \epsilon) \cdot \ell$, $w = O(\log \ell + \lambda/\epsilon)$, then randomly sampled $m \times \ell$ random band matrices have full column rank, time $O(\ell w)$ Gaussian elimination, and time $O(\ell w)$ vector multiplication, with probability $1 - 2^{-\lambda}$

OUR MAIN TOOL: RANDOM-BAND MATRICES [DW19]

- Intuition for full rank:
 - Each band short, so intersects with few other bands
 - But still long enough to find a pivot

1					
0	1	1			
0	1	0			
1	1	0			
1	0	1	1		
0	1	1	1		
	0	0	0	...	
			1		
			0		
			1		

OUR MAIN TOOL: RANDOM-BAND MATRICES [DW19]

- Intuition for fast Gaussian elimination: first sort columns by first non-zero position
 - Each band intersects with few other bands, so each row also consists of a single short band (Chernoff: length $O(w)$)
 - Each row band (by def) intersect with few other bands
 - Thus, getting to echelon form fast
 - Also, back substitution only on $\sim w$ entries for each row

1				
0	1	1		
0	1	0		
1	1	0		
1	0	1	1	
0	1	1	1	
	0	0	0	...
			1	
			0	
			1	

OBLIVIOUS CIPHERTEXT COMPRESSION


Compress: (assume Enc is additively homomorphic)

<i>Enc</i> (20)
<i>Enc</i> (0)
<i>Enc</i> (0)
<i>Enc</i> (14)
<i>Enc</i> (3)
<i>Enc</i> (0)
<i>Enc</i> (35)

c

OBLIVIOUS CIPHERTEXT COMPRESSION

Compress: (assume Enc is additively homomorphic)



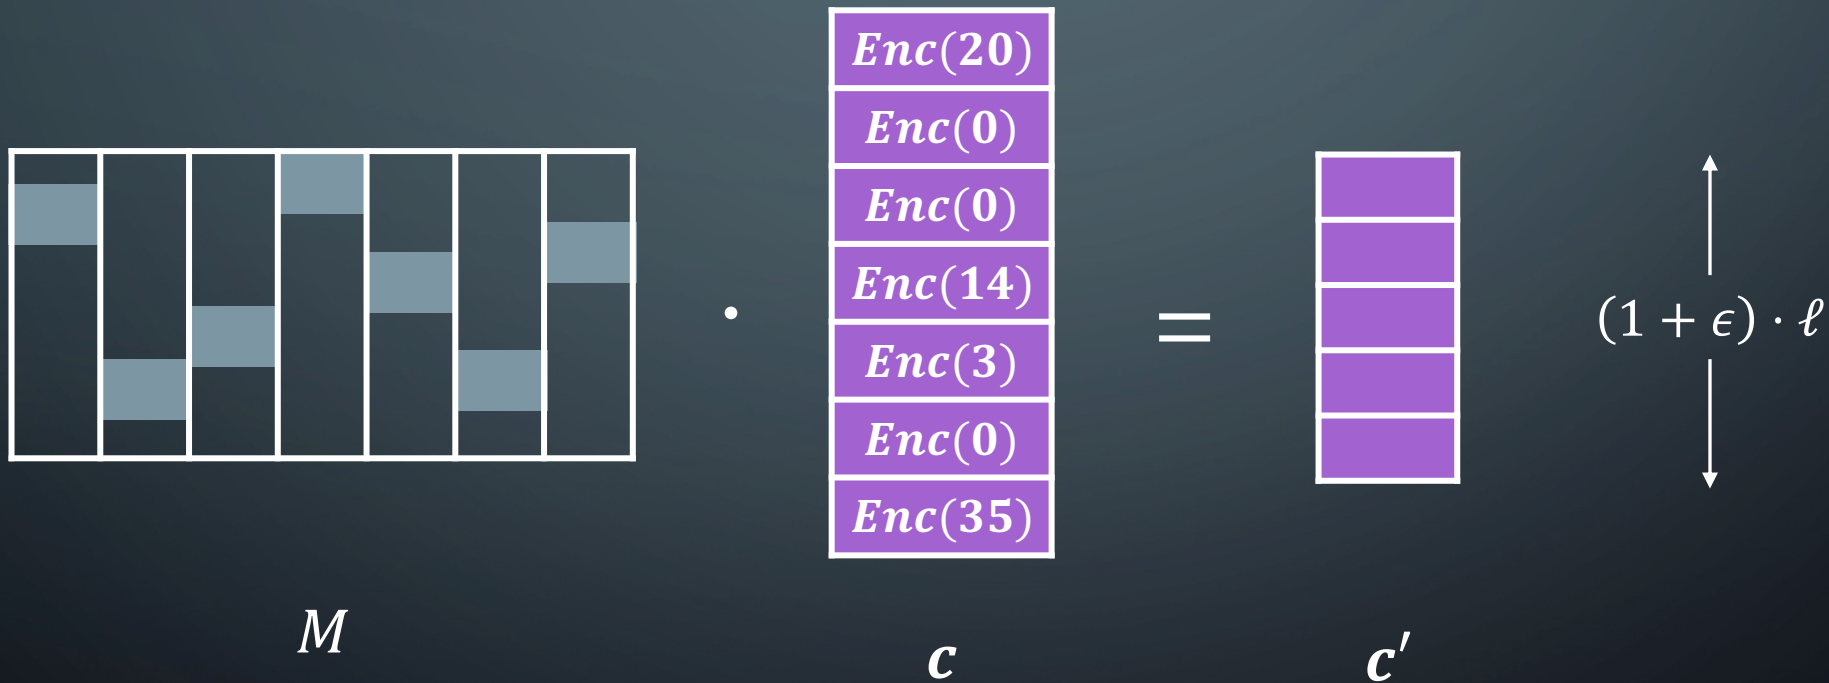
M

$\text{Enc}(20)$
$\text{Enc}(0)$
$\text{Enc}(0)$
$\text{Enc}(14)$
$\text{Enc}(3)$
$\text{Enc}(0)$
$\text{Enc}(35)$

c

OBLIVIOUS CIPHERTEXT COMPRESSION

Compress: (assume Enc is additively homomorphic)



OBLIVIOUS CIPHERTEXT COMPRESSION

Compress: (assume Enc is additively homomorphic)

$$M \cdot \begin{bmatrix} \text{Enc}(20) \\ \text{Enc}(0) \\ \text{Enc}(0) \\ \text{Enc}(14) \\ \text{Enc}(3) \\ \text{Enc}(0) \\ \text{Enc}(35) \end{bmatrix} = \begin{bmatrix} \phantom{\text{Enc}} \\ \phantom{\text{Enc}} \\ \phantom{\text{Enc}} \\ \phantom{\text{Enc}} \\ \phantom{\text{Enc}} \\ \phantom{\text{Enc}} \\ \phantom{\text{Enc}} \end{bmatrix} = \sum_{i=1}^n \mathbf{c}_i \cdot M^i$$

columns

OBLIVIOUS CIPHERTEXT COMPRESSION


Compress: (assume Enc is additively homomorphic)

$$M \cdot \begin{bmatrix} \text{Enc}(20) \\ \text{Enc}(0) \\ \text{Enc}(0) \\ \text{Enc}(14) \\ \text{Enc}(3) \\ \text{Enc}(0) \\ \text{Enc}(35) \end{bmatrix} = \begin{bmatrix} \text{purple} \\ \text{purple} \\ \text{purple} \\ \text{purple} \\ \text{purple} \\ \text{purple} \\ \text{purple} \end{bmatrix} = \sum_{i=1}^n c_i \cdot M^i$$

M c c'

OBLIVIOUS CIPHERTEXT COMPRESSION

Decompress: (client knows nonzero indices)




M



c'

OBLIVIOUS CIPHERTEXT COMPRESSION

Decompress: (client knows nonzero indices)




M



p'

OBLIVIOUS CIPHERTEXT COMPRESSION

Decompress: (client knows nonzero indices)




M'



p'


OBLIVIOUS CIPHERTEXT COMPRESSION

Decompress: (client knows nonzero indices)



OBLIVIOUS CIPHERTEXT COMPRESSION

Decompress: (client knows nonzero indices)



OBLIVIOUS CIPHERTEXT DECOMPRESSION

PIR Requests



ℓ -out-of- n non-zero

- Client compresses and encrypts $\mathbf{m} = (m_1, \dots, m_n)$ to \mathbf{c} of size $\ell + \epsilon$, ϵ small



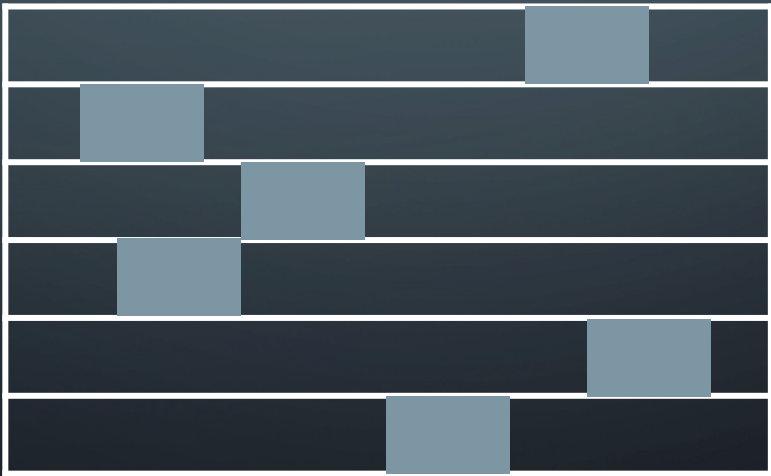
- Server decompresses \mathbf{c} to $\mathbf{c}' = (c_1, \dots, c_n)$ such that if $m_i \neq 0$, c_i encrypts m_i

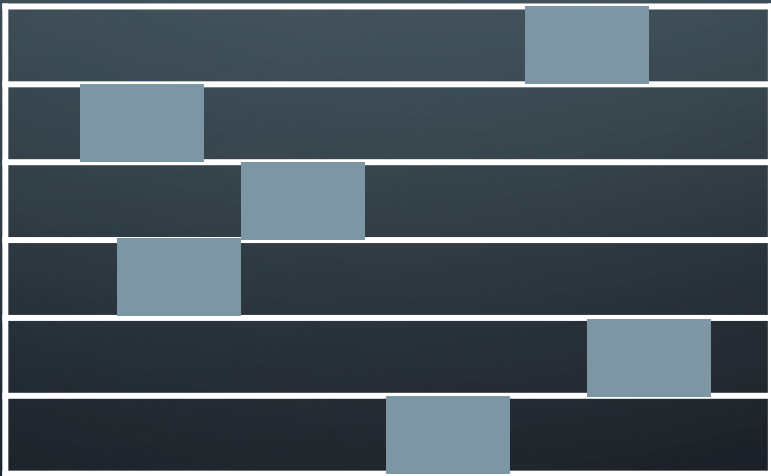



Server must not
learn non-zero
 m_i 's or their
locations

OUR MAIN TOOL: RANDOM-BAND MATRICES [DW19]

- Given $w \leq m$, sample $\ell \times m$ matrix M row-by-row as follows:
 - Pick random starting location from $1, \dots, m - w$
 - Pick random w -bit string
 - Embed random w -bit string at starting location; everything else 0



- # OUR MAIN TOOL: RANDOM-BAND MATRICES [DW19]
- Given $w \leq m$, sample $\ell \times m$ matrix M row-by-row as follows:
 - Pick random starting location from $1, \dots, m - w$
 - Pick random w -bit string
 - Embed random w -bit string at starting location; everything else 0
- 



OUR MAIN TOOL: RANDOM-BAND MATRICES [DW19]

- **Theorem:** There exists some constant $\epsilon > 0$ such that if $m = (1 + \epsilon) \cdot \ell$ and $w = O(\log \ell + \lambda/\epsilon)$, then randomly sampled $\ell \times m$ random band matrices have full row rank, time $O(\ell w)$ Gaussian elimination, and time $O(\ell w)$ vector multiplication, with probability $1 - 2^{-\lambda}$

OBLIVIOUS CIPHERTEXT DECOMPRESSION


Compress (client knows nonzero indices):

20
5
0
0
14
3
0

p

OBLIVIOUS CIPHERTEXT DECOMPRESSION

Compress (client knows nonzero indices):




M

20
5
0
0
14
3
0

p

OBLIVIOUS CIPHERTEXT DECOMPRESSION

Compress (client knows nonzero indices):




M

20
5
0
0
14
3
0

p

OBLIVIOUS CIPHERTEXT DECOMPRESSION

Compress (client knows nonzero indices):




M

20
5
0
0
14
3
0

p


OBLIVIOUS CIPHERTEXT DECOMPRESSION

Compress (client knows nonzero indices):




OBLIVIOUS CIPHERTEXT DECOMPRESSION

Compress (client knows nonzero indices):




OBLIVIOUS CIPHERTEXT DECOMPRESSION

Decompress: (assume Enc is additively homomorphic)




OBLIVIOUS CIPHERTEXT DECOMPRESSION

Decompress: (assume Enc is additively homomorphic)




OBLIVIOUS CIPHERTEXT DECOMPRESSION

Decompress: (assume Enc is additively homomorphic)



OBLIVIOUS CIPHERTEXT DECOMPRESSION

Decompress: (assume Enc is additively homomorphic)



TECHNICALITIES/ADDITIONAL RESULTS

- Cannot simultaneously use Oblivious Ciphertext Decompression and Compression in Batch PIR
 - State of the art Batch PIR uses FHE---accounting for added noise from both techniques makes ciphertexts too big
- Our techniques are compatible with “vectorization” techniques of [MR23]
 - Can fit d requests/responses into a single ciphertext
- (Only) Oblivious Ciphertext Compression also applies to 2-server batch PIR
- From standard Batch PIR + OPRF \Rightarrow Labeled PSI framework, get better Labeled PSI

The image features a dark blue gradient background. In the corners, there are decorative white line art elements resembling circuit boards or neural network connections. These elements consist of thin lines that branch out and terminate in small circles. The top-left and bottom-left corners have more complex, dense patterns, while the top-right and bottom-right corners have simpler, more sparse patterns.

EXPERIMENTS

DB Entry Size	Batch Size & Schemes	Public Param Size	Request Size	Response Size	Total Server Time	Amortized Server Time	Total Client Time	Server Monetary Cost
8 KB	$\ell = 512$							
	Baseline	20.87 MB	1.81 MB	15.59 MB	840 s	1.64 s	6.1 s	\$0.00646
	LSObvCompress	22.62 MB	1.81 MB	10.92 MB	890 s	1.74 s	6.7 s	\$0.00633
	LSObvDecompress	20.87 MB	1.40 MB	15.59 MB	863 s	1.69 s	6.4 s	\$0.00656
	$\ell = 1024$							
	Baseline	20.87 MB	3.35 MB	31.18 MB	1,256 s	1.23 s	6.8 s	\$0.01043
	LSObvCompress	23.11 MB	3.35 MB	21.84 MB	1,369 s	1.34 s	8.2 s	\$0.01025
	LSObvDecompress	20.87 MB	2.55 MB	31.18 MB	1,323 s	1.29 s	8.0 s	\$0.01075
	$\ell = 2048$							
	Baseline	20.87 MB	3.96 MB	62.37 MB	1,750 s	0.85 s	7.0s	\$0.01617
	LSObvCompress	23.43 MB	3.96 MB	43.67 MB	1,871 s	0.91 s	9.7 s	\$0.01520
	LSObvDecompress	20.87 MB	3.15 MB	62.37 MB	1,812 s	0.89 s	9.4 s	\$0.01646
16 KB	$\ell = 512$							
	Baseline	20.87 MB	1.81 MB	31.18 MB	1,286 s	2.51 s	6.3 s	\$0.01047
	LSObvCompress	22.62 MB	1.81 MB	21.84 MB	1,348 s	2.63 s	8.4 s	\$0.00999
	LSObvDecompress	20.87 MB	1.40 MB	31.18 MB	1,308 s	2.55 s	8.3 s	\$0.01056
	$\ell = 1024$							
	Baseline	20.87 MB	3.35 MB	62.37 MB	1,775 s	1.73 s	7.9 s	\$0.01626
	LSObvCompress	23.11 MB	3.35 MB	43.69 MB	1,929 s	1.88 s	9.6 s	\$0.01548
	LSObvDecompress	20.87 MB	2.55 MB	62.37 MB	1,881 s	1.83 s	9.4 s	\$0.01681
	$\ell = 2048$							
	Baseline	20.87 MB	3.96 MB	124.74 MB	2,634 s	1.29 s	8.5 s	\$0.02694
	LSObvCompress	23.43 MB	3.96 MB	87.34 MB	2,773 s	1.35 s	12.4 s	\$0.02439
	LSObvDecompress	20.87 MB	3.15 MB	124.74 MB	2,746 s	1.34 s	12.4 s	\$0.02752

Figure 5: Evaluations of Spiral Batch PIR [14, 58] with and without our compression techniques, LSObvCompress and LSObvDecompress with $\epsilon = 0.05$. We fix the number of entries to $n = 1$ million for all our results.

Batch Size & Schemes	Response Size	Server Time	Client Time	Server Monetary Cost
$\ell = 512$				
Baseline	221 KB	9.63 s	0.01 s	\$0.000076
LSObvCompress	155 KB	9.69 s	0.08 s	\$0.000070
$\ell = 1024$				
Baseline	442 KB	9.76 s	0.01 s	\$0.000096
LSObvCompress	310 KB	9.92 s	0.16 s	\$0.000085
$\ell = 2048$				
Baseline	885 KB	9.79 s	0.01 s	\$0.000136
LSObvCompress	619 KB	10.09 s	0.33 s	\$0.000114
$\ell = 4096$				
Baseline	1,769 KB	9.81 s	0.01 s	\$0.000216
LSObvCompress	1,238 KB	10.53 s	0.78 s	\$0.000172
$\ell = 8192$				
Baseline	3,539 KB	9.82 s	0.01 s	\$0.000375
LSObvCompress	2,477 KB	11.13 s	1.80 s	\$0.000287

Figure 7: Comparison of DPF based two server batch PIR protocol [2] with and without LSObvCompress ($\epsilon = 0.05$). We fix the number of database entries to $n = 1$ million and each entry size to 288 B for all our results.

Label Size & Schemes	Total Online Comm.	Total Online Time	Server Monetary Cost
512 B			
Cong <i>et al.</i> [28]	33.2 MB	169 s	\$0.00397
LSObvCompress	11.4 MB	304 s	\$0.00279
1024 B			
Cong <i>et al.</i> [28]	66.1 MB	331 s	\$0.00787
LSObvCompress	11.6 MB	355 s	\$0.00311
1536 B			
Cong <i>et al.</i> [28]	103.6 MB	535 s	\$0.01244
LSObvCompress	11.9 MB	446 s	\$0.00367

Figure 8: Comparisons of Cong *et al.* [28]’s labeled PSI and our LSObvCompress-based PSI with $\epsilon = 0.05$. We fix the size of the sender’s set to 1 million and the receiver’s set to 512.

An abstract graphic on the left side of the slide, consisting of a network of thin, light-blue lines and small circles, resembling a circuit board or a neural network diagram. The lines and nodes are concentrated on the left edge and spread out towards the center.

THANKS!