

# Single-Server Private Information Retrieval and Secure Aggregation in the Shuffle Model

Adrià Gascón

Yuval Ishai

Mahimna Kelkar

Baiyu Li


[Yiping Ma](#)

Mariana Raykova




# The two worlds of PIR

Single-server PIR



(Non-colluding) Multi-server PIR



# The two worlds of PIR

## Single-server PIR

### Plain model

Early works: [KO97], [CMS99], [GR05], ...

Recent: XPIR, SEAL-PIR, Spiral, WhisPIR, ...

DCR, DDH, ...

(R)LWE

Making  
the server  
fast

### Preprocessing

[PPY18], SimplePIR, FrodoPIR, Piano, HintlessPIR, ...

[BIM04]

## (Non-colluding) Multi-server PIR

### Plain model

Early works: from coding theory, e.g., [CGKS95]

Recent: DPF-based, efficient construction from [BGI16]

Deployment cost

Concretely fast

### Preprocessing

[BIM04], [CK20], [KC21], [ISW24], ...

# The two worlds of PIR

## Single-server PIR

### Plain model

Early works: [KO97], [CMS99], [GR05], ...  
Recent: XPIR, SEAL-PIR, Spiral, WhisPIR, ...

DCR, DDH, ...

(R)LWE

Making  
the server  
fast

### Preprocessing

[PPY18], SimplePIR, FrodoPIR, Piano, HintlessPIR, ...

[BIM04]

## (Non-colluding) Multi-server PIR

### Plain model

Early works: from coding theory, e.g., [CGKS95]  
Recent: DPF-based, efficient construction from [BGI16]

Deployment cost

Concretely fast



Unfriendly to database updates, big client storage, still fall behind the fast multi-server PIR, ...

*Can we have a single-server PIR that is as fast as multi-server PIR?*



# Best of both worlds?

- Revisiting the shuffle model gives us ideas
  - Common in differential privacy literature
  - Proposed in [IKOS06] in the context of PIR (but later didn't receive much attention)
- The shuffle model assumption
  - Many clients make queries simultaneously
  - The queries are shuffled before reaching the server

A two-way  
channel



# Best of both worlds?

- Revisiting the shuffle model gives us ideas
  - Common in differential privacy literature
  - Proposed in [IKOS06] in the context of PIR (but later didn't receive much attention)
- The shuffle model assumption
  - Many clients make queries simultaneously  Many applications naturally has the feature
  - The queries are shuffled before reaching the server  Studied in many works on anonymity

# Best of both worlds? Yes, in the shuffle model

- Prior results:
  - [IKOS06]: a construction based on non-standard computational assumption
  - [IKLM24]: a class of constructions for IT-secure PIR with sublinear communication

Both works are not concretely efficient


- This work: single-server PIR in the shuffle model that matches the (server) cost of a very fast multi-server PIR

# This work


- A construction for single-server PIR in the shuffle model
  - Based on LPN and conjectured MDSD hardness
  - 9–25X improvement on throughput over SimplePIR [HHCG<sup>+</sup>], depending on parameters
- Computationally secure aggregation in the shuffle model
  - Same assumptions as the PIR construction
  - 25X savings for communication compared to existing best statistical scheme [BBGN20]



# Outline

- 
- Shuffle PIR: the security goal
  - The “split-and-mix” paradigm
  - Our constructions
    - Shuffle PIR based on LPN, MDSD
    - By-product: computationally secure aggregation in the shuffle model
  - Discussion and open questions



# PIR in the shuffle model: the security goal



**Anonymity does not imply message privacy:**  
It hides who sends what,  
but does not hide the messages themselves



# PIR in the shuffle model: the security goal

- Anonymization does not trivialize the PIR problem!




# PIR in the shuffle model: the security goal

- Anonymization does not trivialize the PIR problem!




# Outline

- Shuffle PIR: the security goal
-  • The “split-and-mix” paradigm
- Our constructions
  - Shuffle PIR based on LPN, MDSD
  - By-product: computationally secure aggregation in the shuffle model
- Discussion and open questions

# The **split-and-mix** paradigm

- Privacy from anonymity [IKOS06]: computing sum




Take a large enough  $p$ , each client splits its inputs into  $k$  shares in  $\mathbb{Z}_p$



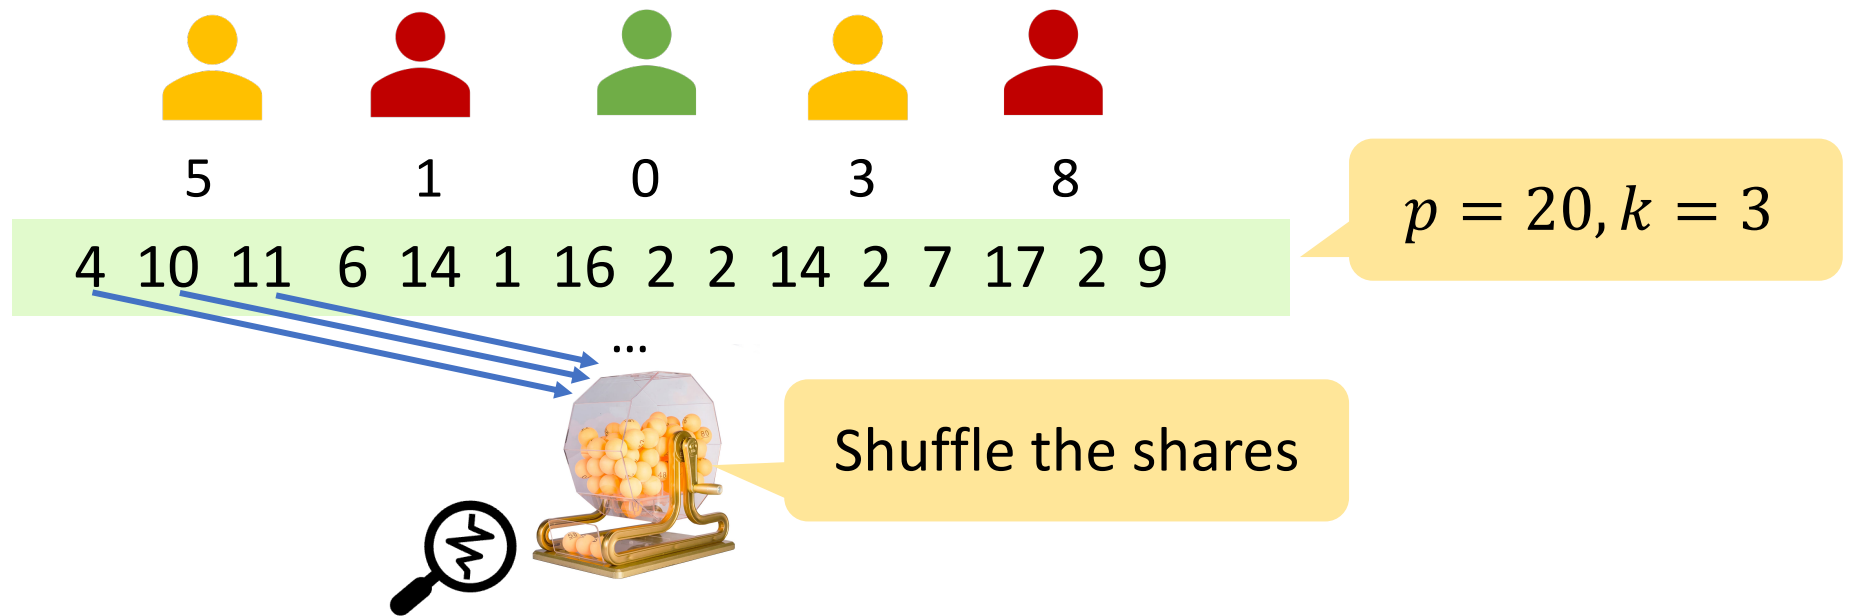
# The **split-and-mix** paradigm

- Privacy from anonymity [IKOS06]: computing sum



# The **split-and-mix** paradigm

- Privacy from anonymity [IKOS06]: computing sum





The shuffled shares leak nothing about the inputs except the sum



# The **split-and-mix** paradigm

- Privacy from anonymity [IKOS06]: computing sum



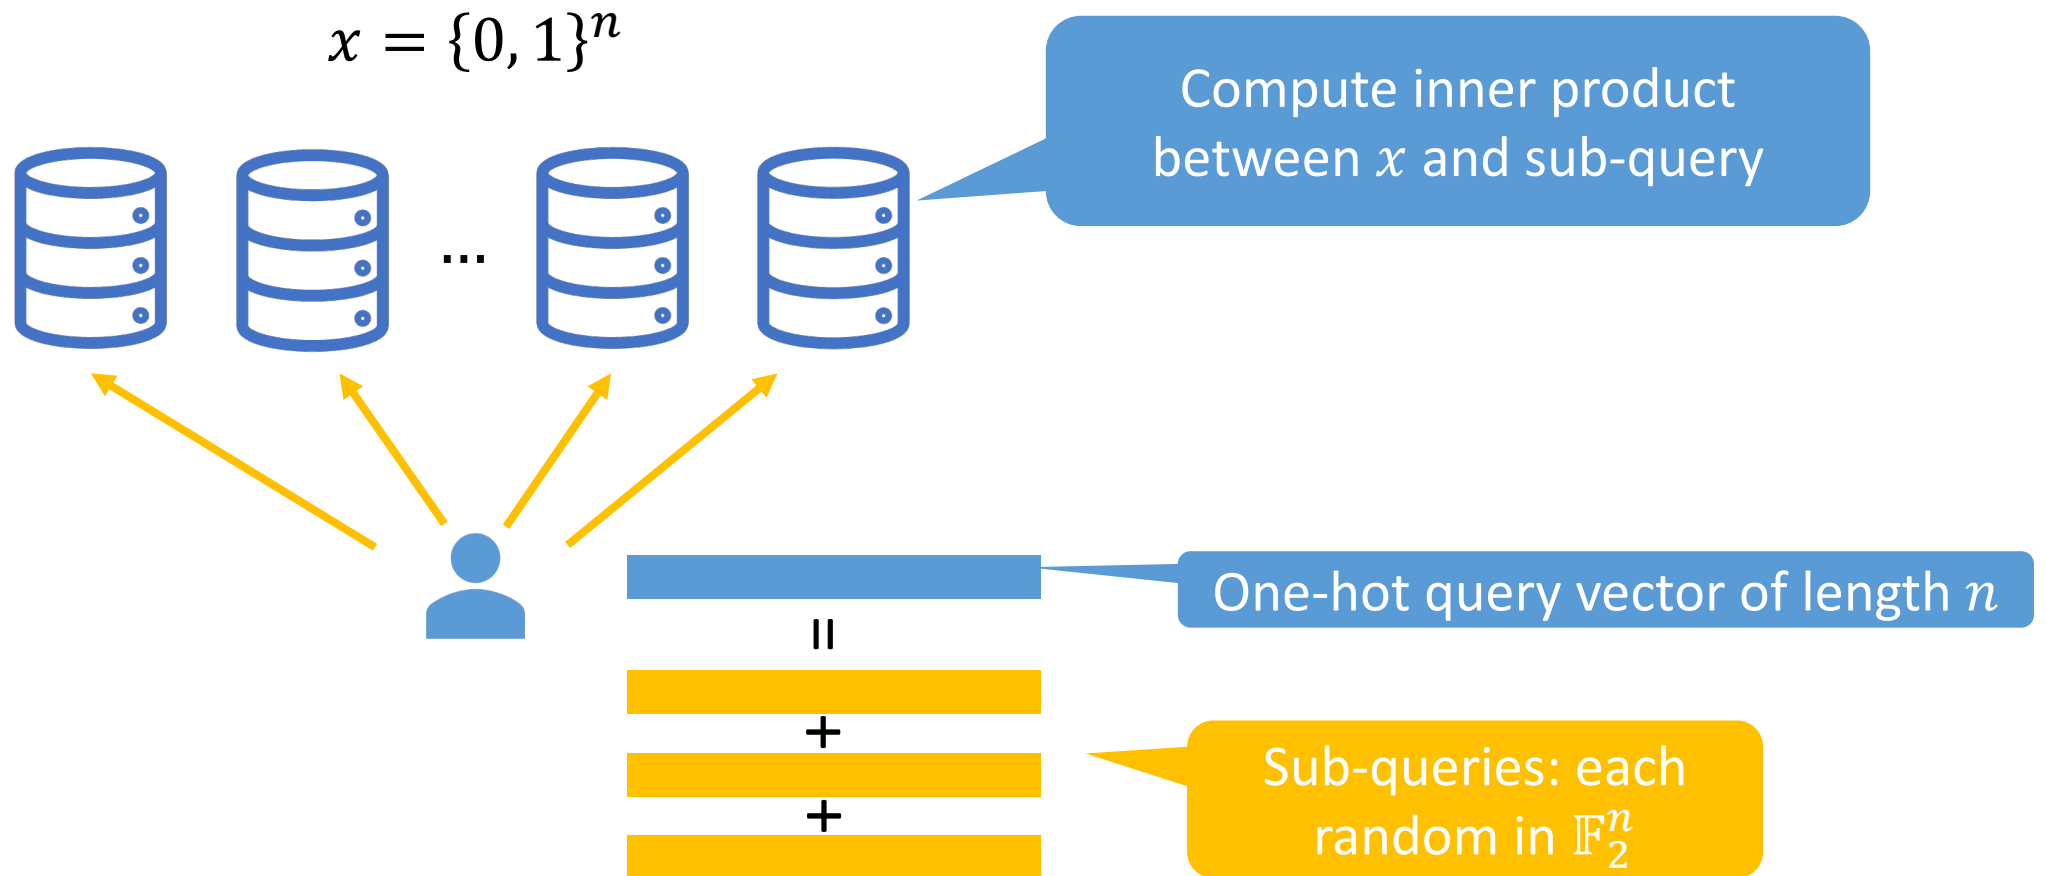
Any two different configurations with equal sum

$$\text{View}(10, 2, 2, 1, 1) \approx \text{View}(4, 4, 4, 4, 0)$$

The equation is displayed on a yellow background. Above the equation is an icon of a lottery ball machine with a magnifying glass over a Bitcoin symbol.


Statistical/computational

# Split-and-mix for PIR



# Split-and-mix for PIR



$$x = \{0, 1\}^n$$



Mixing sub-queries?

# Split-and-mix for PIR

“PIR in the secret sharing paradigm”: view the sub-queries as shares, then mix the sub-queries




# Split-and-mix for PIR

“PIR in the secret sharing paradigm”: view the sub-queries as shares, then mix the sub-queries




Any two different set of query indices with equal “sum”

Each input is split to  $k$  sub-queries



$View(10, 2, 2, 1, 1) \approx View(4, 4, 4, 4, 0)$



This work: computational


# Outline

- Shuffle PIR: the security goal
- The “split-and-mix” paradigm
- Our constructions
  - Shuffle PIR based on LPN, MDSD
  - By-product: computationally secure aggregation in the shuffle model
- Discussion and open questions




# Our construction: security analysis

- Relevant computational assumptions:  $k$ -sum, syndrome decoding (dual-LPN)



# Our construction: security analysis

- Let's do a slight modification...





# Our construction: security analysis

- Decisional syndrome decoding (dual LPN) [BFKL94, AIK07]

A random matrix

A low-weight  
random vector

A random  
vector

$$(H, H \cdot e) \approx_c (H, r)$$

Everything in  $\mathbb{F}_2$


$$( \text{Matrix with stripes} , \text{Orange vector} ) \approx_c ( \text{Solid blue matrix} , \text{Blue vector} )$$

# Reducing to LPN

- Considering a single client

- $H \cdot e + y = q,$

where  $q \in \mathbb{F}_2^n, H \in \mathbb{F}_2^{n \times m}$  and  $e \in \mathbb{F}_2^m$  with Hamming weight  $k$



# Conjectured MDSD security

- When there are  $c$  clients
- $H \cdot E + Y = Q$ , where  $E$  is a “structured” matrix, and  $Q \in \mathbb{F}_2^{n \times c}$

## Multi-Disjoint Syndrome Decoding (MDSD)

A random matrix

A structured  
matrix with low-  
weight columns

A random  
matrix

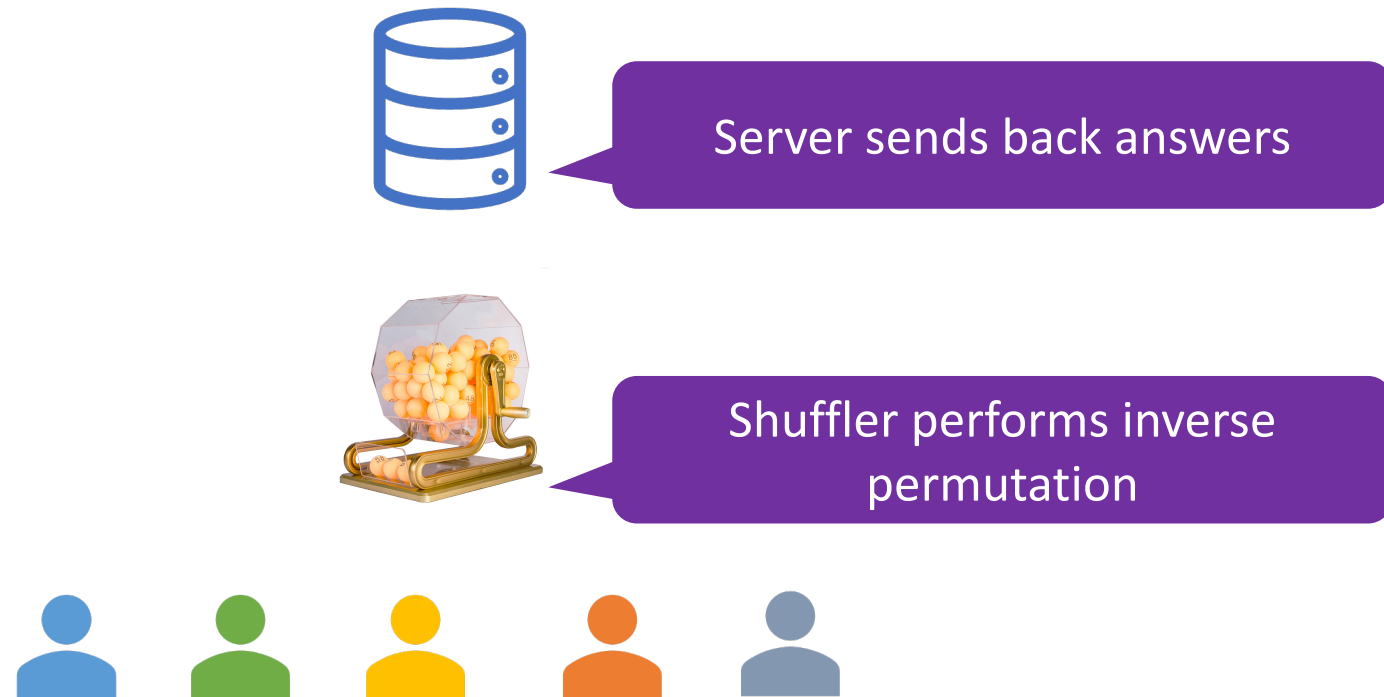
$$(H, H \cdot E) \approx_c (H, R)$$

$H$  has shares and dummies

# Shuffle PIR: the computational construction




# Shuffle PIR: the computational construction




# Optimization for concrete efficiency

- PRG compression
- Batch processing
- Parameter slicing
- Offline/online model

# Optimization: PRG compression



# Optimization: batch processing



- ❖ Random access to a large chunk is expensive
- ❖ Sequential access is much faster than random access



# Optimization: batch processing

Inner product = sum of the “selected” entries

Database


Share vector 1	(1 0 ...	0 1 0 1 ...	0 1 1 ...	0 1 0 ...	1 0 )
Share vector 2	(0 0 ...	0 0 1 1 ...	1 1 1 ...	1 0 0 ...	0 0 )
Share vector 3	(1 1 ...	0 1 0 1 ...	0 1 1 ...	0 1 0 ...	1 1 )
Share vector 4	(0 0 ...	0 1 1 1 ...	0 0 1 ...	0 1 1 ...	0 1 )

# Optimization: parameter slicing

- The issue of too many dummies:
  - For a database of size  $\sim 1K$ , if each query vector is split into 10 shares, then we need in total  $\sim 1M$  dummies
  - Many dummies  $\Rightarrow$  high (and wasteful) anonymity cost and server computation

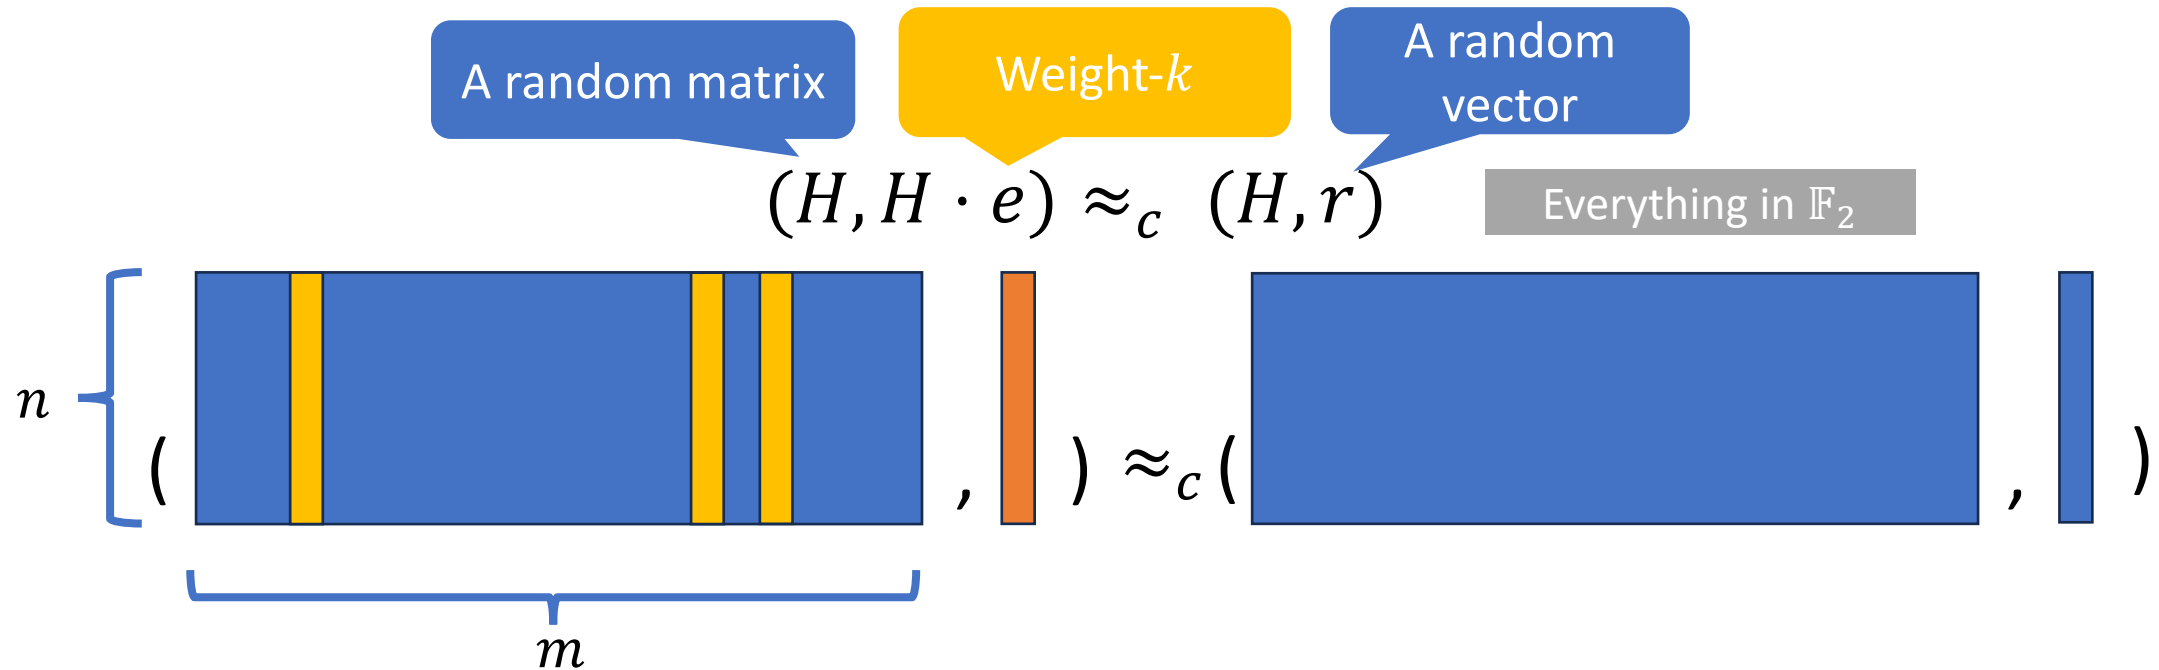
# Optimization: parameter slicing

- Let the SD problem be parameterized by  $(m, n, k)$
- Fix  $m, k$ , the smaller  $n$  is, the harder the SD instance is




# Optimization: parameter slicing

- Let the SD problem be parameterized by  $(m, n, k)$
- Fix  $m, k$ , the smaller  $n$  is, the harder the SD instance is



# Optimization: parameter slicing

- Let the SD problem be parameterized by  $(m, n, k)$
- Fix  $m, k$ , the smaller  $n$  is, the harder the SD problem is



Dimension  $n' \times m$

Weight  $k$

$$(H, H \cdot e) \approx_c (H, r)$$


Everything in  $\mathbb{F}_2$

# Optimization: database slicing

- Let the SD problem be parameterized by  $(m, n, k)$
- Fix  $m, k$ , the smaller  $n$  is, the harder the SD problem is
- A breakeven point for  $k$  and  $n'$ : when  $n'$  becomes smaller,
  - 😊 • The corresponding MDSD instance is harder
  - 😞 • More MDSD instances, and we need to guarantee the adversary cannot break any of these MDSD instances (can be analyzed by union bound)

# Optimization: offline/online model

- Offline: client sends  $k - 1$  random vectors to the server for preprocessing
- Online: set the last share = query vector – the random vectors



Now we don't shuffle all of them together—will it affect security?

# Performance

- PRG compression, batch processing, database slicing, offline/online

Database	# Clients	Our PIR Protocol in Construction 3				SimplePIR		HintlessPIR (ms)	Spiral (ms)
		Offline latency (ms)	Online latency (ms)	Online throughput (GiB/s)	Online throughput (batched, GiB/s)	Latency (ms)	Throughput (GiB/s)		
$2^{20} \times 256$ bytes	1K	132	7	36.64	125.05	45	5.56	575	794
	10K	63							
	100K	47							



# Outline

- Shuffle PIR: the security goal
- The “split-and-mix” paradigm
- Our constructions
  - Shuffle PIR based on LPN, MDSD
  - By-product: computationally secure aggregation in the shuffle model
- Discussion and open questions




# Secure aggregation based on MDSD

Isn't the same with the  
PIR construction?



Dummies will change the sum!




# Secure aggregation based on MDSD

What should we do?



Using the shuffler to  
insert shares of zeros!



# Long-vector aggregation: optimization

- Same ideas used in many places: key-and-message homomorphic encryption

Short Long

$$Enc(s_1, x_1) + Enc(s_2, x_2) = Enc(s_1 + s_2, x_1 + x_2)$$

We use RLWE to compactly pack data


- $c$  clients send  $Enc(s_1, x_1), \dots, Enc(s_c, x_c)$
- Use the shuffle scheme for aggregating  $s_1, \dots, s_c$   
(Similar as before, we can apply PRG compression)

# Performance

Input $\times$ Field	Input size (KiB)	# Clients	Aggregation Protocol in Construction 1			Info-Theoretic Protocol [BBGN20]		
			# Shares	Upload size (KiB)	Expansion ratio	# Shares	Upload size (KiB)	Expansion ratio
$2^{15} \times \mathbb{F}_2$	4	100	405	10	2.57	6317	103	25.67
		1000	88	5	1.34	3856	64	16.06
		10000	37	5	1.14	2775	47	11.84
$2^{15} \times \mathbb{F}_{65537}$	64	100	410	70	1.10	100819	1639	25.61
		1000	77	65	1.02	61525	1025	16.02
		10000	33	65	1.01	44271	756	11.81

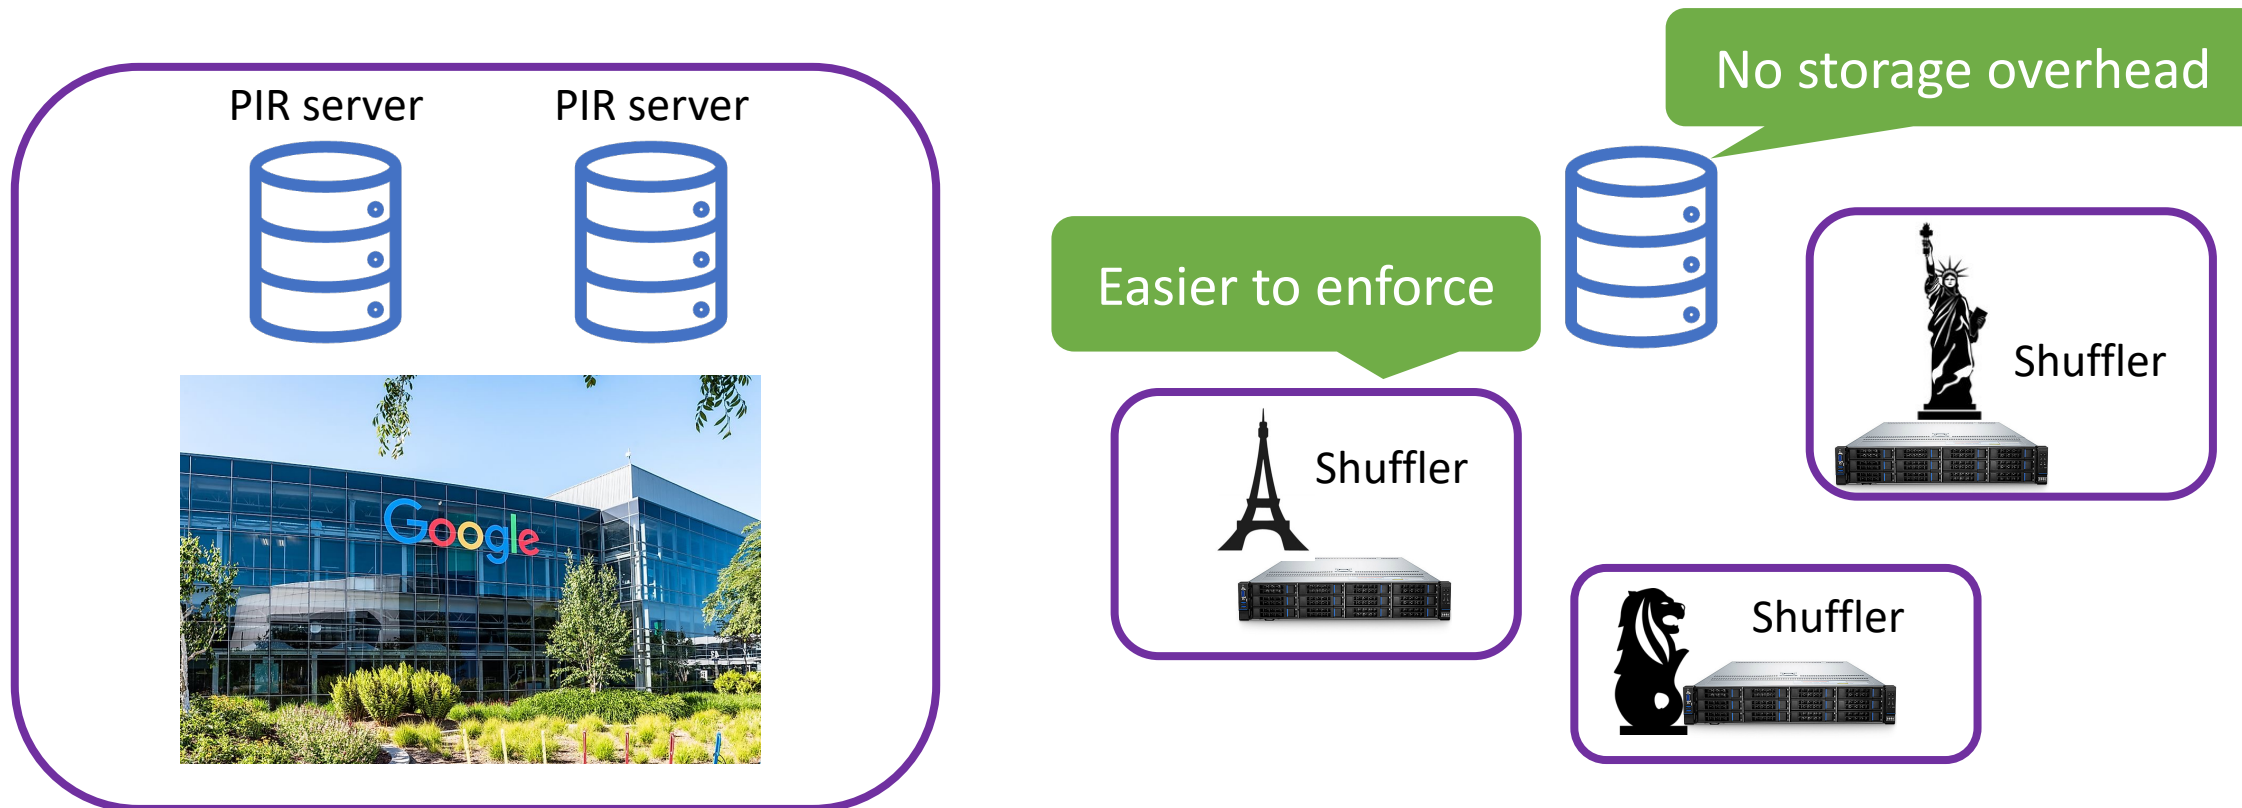
Expansion ratio measures the communication overhead of the protocol over the input vector size

# Outline

- Shuffle PIR: the security goal
- The “split-and-mix” paradigm
- Our constructions
  - Shuffle PIR based on LPN, MDSD
  - By-product: computationally secure aggregation in the shuffle model
-  • Discussion and open questions

# Discussion

- Assuming non-colluding servers vs. assuming a two-way anonymous channel



# Discussion

- We are in the situation of exploiting tradeoffs: making assumptions, altering models (different types of preprocessing, relaxed security, etc.)
- Guaranteeing different assumptions does not require the same amount of efforts: system efforts, law efforts, etc.
- The likelihood of assumptions being compromised in real-world scenarios may vary



*We should keep these in mind when introducing new models!*



# Backup Slides

# PIR with variable-sized records

- To deploy PIR in real-world applications...



To retrieve privately, it is necessary to hide record size


# PIR with variable-sized records

- To deploy PIR in real-world applications...

Often assume the same size, mostly  $\{0, 1\}^n$



Database entries of PIR in theory




Database records in practice

To retrieve privately, it is necessary to hide record size

# PIR with variable-sized records

- Padding solves the problem: how about efficiency?



Database records in practice


## Features

The discrepancy between the smallest and the largest record can be huge  
Majority of the records are small  
Most users access the small records much more often than the large records

# PIR with variable-sized records

- Padding solves the problem: how about e

Waste of server storage  
(though can virtually store)



## Features

The discrepancy

Majority of the records are small

Most users access the small records much more often than the large records

Client who retrieves the small record has to  
pay the cost of retrieving the largest record

practice


can be huge

# PIR with variable-sized records

- In the “standard” model, there is no way out
- In the shuffle model: yes, we can
  - No server storage overhead
  - Client communication proportional to the length of the retrieved record
  - Leak only the total size of all queried records


# PIR with variable-sized records

- A toy protocol



# PIR with variable-sized records

- A toy protocol

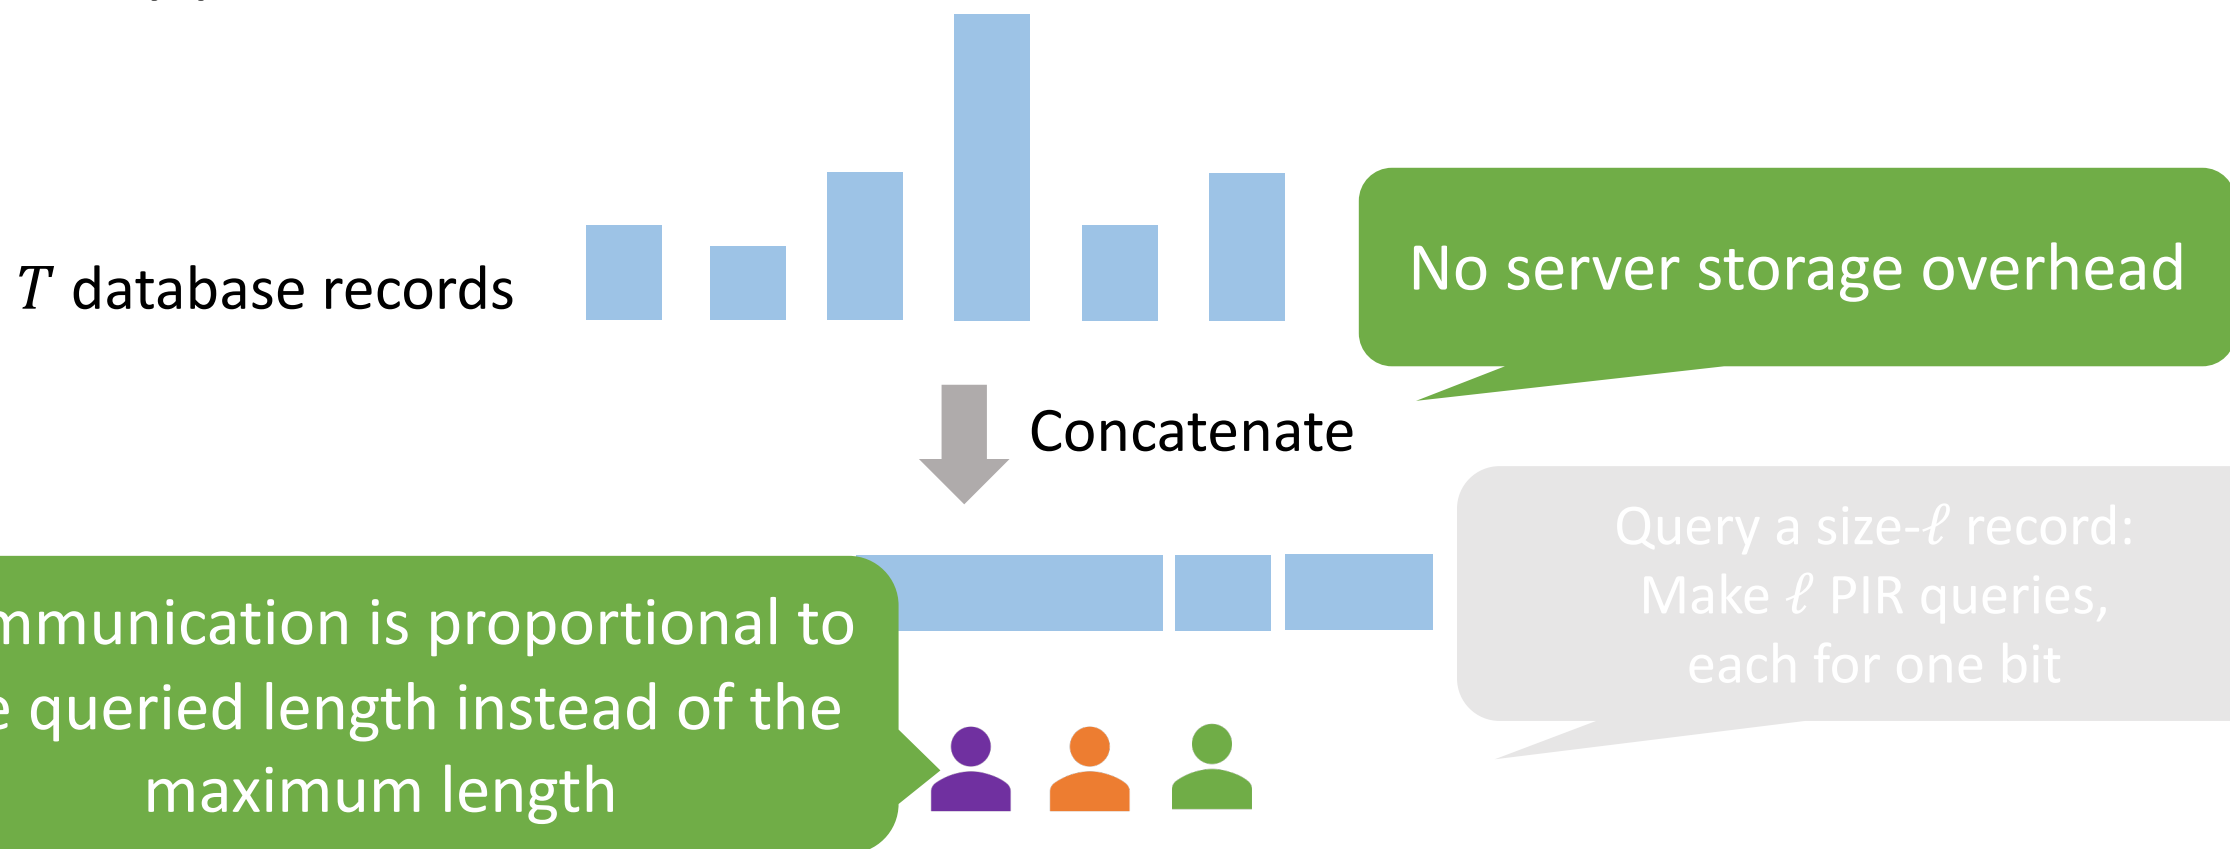


Query a size- $\ell$  record:  
Make  $\ell$  PIR queries,  
each for one bit



# PIR with variable-sized records

- A toy protocol




# PIR with variable-sized records

- A toy protocol



# PIR with variable-sized records

- Revisit the toy protocol




Why not retrieve more bits in each PIR query?

Query a size- $\ell$  record:  
Make  $\ell$  PIR queries,  
each for one bit

# PIR with variable-sized records

- Splitting records to the powers of two



Secure or not?

Deterministic splitting is not secure  
(unless split down to 1)

Server (logically) prepare  $\log n$  databases:  
the  $j$ -th database is partitioned to  $2^j$  bits per entry


# PIR with variable-sized records

- Splitting records to the powers of two

Consider 5 1 1 1


v.s.

2 2 2 2




# PIR with variable-sized records

- Our approach: recursive splitting




# PIR with variable-sized records

- Our approach: recursive splitting



# PIR with variable-sized records


- Our approach: recursive splitting






# PIR with variable-sized records

- Our approach: recursive splitting



# PIR with variable-sized records

- Our approach: recursive splitting



The final blocks that the client will retrieve (using PIR)

# PIR with variable-sized records


- A complication of recursive splitting: fully split the highest  $\log C$  levels

Consider 5 1 1 1

v.s.

2 2 2 2

With 1/2 probability, there will be a block



# PIR with variable-sized records


- A complication of recursive splitting: fully split the highest  $\log C$  levels

Consider  $M-3 \quad 1 \quad 1 \quad 1$

v.s.

$M/4 \quad M/4 \quad M/4 \quad M/4$

With 1/2 probability, there will be a block




# PIR with variable-sized records

- A complication of recursive splitting: fully split the highest  $\log C$  levels

Consider  $M-3$  1 1 1

v.s.



$M/4$   $M/4$   $M/4$   $M/4$



As long as there are sufficient number  
of blocks at this level

# PIR with variable-sized records

- Splitting records to the power of two



The multi-set of record lengths from all clients will not leak any individual queried length